

Deep Learning for NLP

Basics


Word Embeddings



Male-Female



Verb Tense



Country-Capital

Learning goals

- Understand what word embeddings are
- Learn the main methods for creating them

MOTIVATION (1) I

- How to represent words/tokens in a neural network?
- Possible solution: one-hot encoded indicator vectors of length $|V|$.

$$\vec{w}^{(\text{the})} = \begin{bmatrix} 1 \\ 0 \\ 0 \\ \vdots \end{bmatrix} \quad \vec{w}^{(\text{cat})} = \begin{bmatrix} 0 \\ 1 \\ 0 \\ \vdots \end{bmatrix} \quad \vec{w}^{(\text{dog})} = \begin{bmatrix} 0 \\ 0 \\ 1 \\ \vdots \end{bmatrix}$$

- **Question:** Why is this a bad idea?
 - Parameter explosion ($|V|$ might be $> 1M$)
 - All word vectors are orthogonal to each other
→ no notion of word similarity

MOTIVATION (2)

- Learn one word vector $\vec{w}^{(i)} \in \mathbb{R}^D$ (“word embedding”) per word i
- Typical dimensionality: $50 \leq D \leq 1000 \ll |V|$
- Embedding matrix: $\mathbf{W} \in \mathbb{R}^{|V| \times D}$
- **Question:** Advantages of using word vectors?

MOTIVATION (2)

- Learn one word vector $\vec{w}^{(i)} \in \mathbb{R}^D$ (“word embedding”) per word i
- Typical dimensionality: $50 \leq D \leq 1000 \ll |V|$
- Embedding matrix: $\mathbf{W} \in \mathbb{R}^{|V| \times D}$
- **Question:** Advantages of using word vectors?
 - We can express similarities between words, e.g., with cosine similarity:

$$\cos(\vec{w}^{(i)}, \vec{w}^{(j)}) = \frac{\vec{w}^{(i)T} \vec{w}^{(j)}}{\|\vec{w}^{(i)}\|_2 \cdot \|\vec{w}^{(j)}\|_2}$$

- Since the embedding operation is a *lookup operation*, we only need to update the vectors that occur in a given training batch

MOTIVATION (3) I

Supervised training?

- *Training embeddings from scratch:*
 - Initialize randomly and learn it during training phase
 - Words that play similar roles w.r.t. task get similar embeddings
- *Example: Sentiment Classification*
 - We might expect $\vec{w}^{(\text{great})} \approx \vec{w}^{(\text{awesome})}$
- **Question:** What could be a problem at test time?
 - If training set is small, many words are unseen during training and therefore have random vectors
- We typically have more unlabeled than labeled data.
 - Can we learn embeddings from the unlabeled data?

MOTIVATION (4) I

- Distributional hypothesis:
“A word is characterized by the company it keeps” (J.R. Firth, 1957)
- *Idea:*
Learn similar vectors for words that occur in similar contexts
- Three different (milestone) methods:
 - Word2Vec ▶ (Mikolov et al., 2013)
 - GloVe (not covered) ▶ (Pennington et al., 2014)
 - FastText ▶ (Bojanowski et al., 2016)

WORD2VEC AS A BIGRAM LANGUAGE MODEL I

Model architecture:

- Words in our vocabulary are represented as two sets of vectors:
 - $\vec{w}^{(i)} \in \mathbb{R}^D$ if they are to be predicted
 - $\vec{v}^{(i)} \in \mathbb{R}^D$ if they are conditioned on as context
- Predict word i given previous word j :

$$P(i|j) = f(\vec{w}^{(i)}, \vec{v}^{(j)})$$

- **Question:** What is a possible function $f(\cdot)$?

WORD2VEC AS A BIGRAM LANGUAGE MODEL

- Softmax!

$$P(i|j) = \frac{\exp(\vec{w}^{(i)T} \vec{v}^{(j)})}{\sum_{k=1}^{|V|} \exp(\vec{w}^{(k)T} \vec{v}^{(j)})}$$

- **Question:** Problem with training softmax?

WORD2VEC AS A BIGRAM LANGUAGE MODEL

- Softmax!

$$P(i|j) = \frac{\exp(\vec{w}^{(i)T} \vec{v}^{(j)})}{\sum_{k=1}^{|V|} \exp(\vec{w}^{(k)T} \vec{v}^{(j)})}$$

- **Question:** Problem with training softmax?

Needs to compute dot products with the whole vocabulary in the denominator for every single prediction

→ *SLOW*

SPEEDING UP TRAINING: NEGATIVE SAMPLING I

- One option: **Hierarchical Softmax** (not covered) reduces complexity from $\mathcal{O}(|V|)$ to $\mathcal{O}(\log_2 |V|)$
- Another trick: **Negative Sampling**
(a variant of *noise contrastive estimation*)
→ Changes the objective function; the resulting model is not a language model anymore!
- **Idea**: Instead of predicting the probability distribution over the whole vocabulary, make binary decisions for a small number of words.
 - “Positive” samples: Bigrams seen in the corpus.
 - “Negative” samples: Random bigrams (not seen in corpus)

NEGATIVE SAMPLING: LIKELIHOOD I

- Given:
 - Positive training set: $\text{pos}(\mathcal{O})$
 - Negative training set: $\text{neg}(\mathcal{O})$

$$L = \prod_{(i,j) \in \text{pos}(\mathcal{O})} P(\text{pos} | \vec{w}^{(i)}, \vec{v}^{(j)}) \prod_{(i',j') \in \text{neg}(\mathcal{O})} P(\text{neg} | \vec{w}^{(i')}, \vec{v}^{(j')})$$

- $P(\text{pos} | \vec{w}, \vec{v}) = \sigma(\vec{w}^T \vec{v})$
- $P(\text{neg} | \vec{w}, \vec{v}) = 1 - P(\text{pos} | \vec{w}, \vec{v})$
- Question:** Why not just maximize $\prod_{(i,j) \in \text{pos}(\mathcal{O})} P(\text{pos} | \vec{w}^{(i)}, \vec{v}^{(j)})$?
 - Trivial solution: make all $\vec{w}, \vec{v}$ identical

WORD2VEC WITH NEGATIVE SAMPLING I

- Maximize likelihood of training data:

$$\mathcal{L}(\theta) = \prod_i P(y^{(i)} | x^{(i)}; \theta)$$

$\Leftrightarrow$ minimize negative log likelihood:

$$NLL(\theta) = -\log \mathcal{L}(\theta) = -\sum_i \log P(y^{(i)} | x^{(i)}; \theta)$$

WORD2VEC WITH NEGATIVE SAMPLING I

- **Question:** What do these components stand for in Word2Vec with negative sampling?
 - $x^{(i)}$ Word pair, from corpus OR randomly created
 - $y^{(i)}$ Label:
 - 1 = word pair is from positive training set,
 - 0 = word pair is from negative training set
 - Parameters θ of the model: $\vec{v}$, $\vec{w}$
 - $P(\dots)$ Logistic sigmoid: $P(1|\cdot) = \sigma(\vec{w}^T \vec{v})$, resp. $P(0|\cdot) = 1 - \sigma(\vec{w}^T \vec{v})$.

SPEEDING UP TRAINING: NEGATIVE SAMPLING I

- Constructing a good negative training set can be difficult
- Often it is some random perturbation of the training data (e.g. replacing the second word of each bigram by a random word).
- The number of negative samples is often a multiple (1x to 20x) of the number of positive samples
- Negative sets are often constructed per batch

QUESTIONS I

- **Question:** How many dot products do we need to calculate for a given word pair? How does this compare to the naive and hierarchical softmax?
 - $M + 1 \approx \log_2 |V| \ll |V|$
(for $M = 20, |V| = 1,000,000$)

SKIP-GRAM (WORD2VEC) I

Create a fake task:

- **Training objective:** Given a word, predict the neighbouring words
- **Generation of samples:** Sliding fixed-size window over the text

Example: Window size = 2

The	quick	brown	fox	jumps	over	the	lazy	dog
-----	-------	-------	-----	-------	------	-----	------	-----

⇒ (the, quick); (the, brown)

The	quick	brown	fox	jumps	over	the	lazy	dog
-----	-------	-------	-----	-------	------	-----	------	-----


⇒ (quick, the); (quick, brown); (quick, fox)

The	quick	brown	fox	jumps	over	the	lazy	dog
-----	-------	-------	-----	-------	------	-----	------	-----

⇒ (brown, the); (brown, quick), (brown, fox), (brown, jumps)

SKIP-GRAM (WORD2VEC) I

- Idea: Learn many bigram language models at the same time.
- Given word $w_{[t]}$, predict words inside a window around $w_{[t]}$:
 - One position before the target word: $p(w_{[t-1]}|w_{[t]})$
 - One position after the target word: $p(w_{[t+1]}|w_{[t]})$
 - Two positions before the target word: $p(w_{[t-2]}|w_{[t]})$
 - ... up to a specified window size c .
- Models share all $\vec{w}$, $\vec{v}$ parameters!



Skip-gram

SKIP-GRAM: OBJECTIVE I

- Optimize the joint likelihood of the $2c$ language models:

$$p(w_{[t-c]} \dots w_{[t-1]} w_{[t+1]} \dots w_{[t+c]} | w_{[t]}) = \prod_{\substack{i \in \{-c \dots c\} \\ i \neq 0}} p(w_{[t+i]} | w_{[t]})$$

- Negative Log-likelihood for whole corpus (of size N):

$$NLL = - \sum_{t=1}^N \sum_{\substack{i \in \{-c \dots c\} \\ i \neq 0}} \log p(w_{[t+i]} | w_{[t]})$$

- Using negative sampling as approximation:

$$\approx - \sum_{t=1}^N \sum_{\substack{i \in \{-c \dots c\} \\ i \neq 0}} \left[\log \sigma(\vec{w}_{[t+i]}^T \vec{v}_{[t]}) + \sum_{m=1}^M \log[1 - \sigma(\vec{w}^{(*)T} \vec{v}_{[t]})] \right]$$

- $\vec{w}^{(*)}$ is the word vector of a random word, M is the number of negatives per positive sample

CONTINUOUS BAG OF WORDS (CBOW) I

Create a fake task:

- **Training objective:** Given a context, predict the center word
- **Generation of samples:** Sliding fixed-size window over the text

Example: Window size = 2

The	quick	brown	fox	jumps	over	the	lazy	dog
-----	-------	-------	-----	-------	------	-----	------	-----

⇒ ([the, quick, fox, jumps], brown)

The	quick	brown	fox	jumps	over	the	lazy	dog
-----	-------	-------	-----	-------	------	-----	------	-----

⇒ ([quick, brown, jumps, over], fox)

The	quick	brown	fox	jumps	over	the	lazy	dog
-----	-------	-------	-----	-------	------	-----	------	-----


⇒ ([brown, fox, over, the], jumps)

CONTINUOUS BAG OF WORDS (CBOW) I

- Like Skipgram, but...
- Predict word $w_{[t]}$, given the words inside the window around $w_{[t]}$:

$$p(w_{[t]} | w_{[t-c]} \dots w_{[t-1]} w_{[t+1]} \dots w_{[t+c]})$$

$$\propto \vec{w}_{[t]}^T \sum_{\substack{i \in -c \dots c \\ i \neq 0}} \vec{v}_{[t+i]}$$



FASTTEXT (1) I

Accomplishments:

- Words can be represented as dense, low-dimensional vectors ✓
- Easy to capture similarity between words ✓
- Additive Compositionality of word vectors ✓

Open issues:

- Even if we train Word2Vec on a very large corpus, we will still encounter unknown words at test time
- What about rare words?
- Orthography can often help us:
- $\vec{w}^{(\text{remuneration})}$ should be similar to
 - $\vec{w}^{(\text{remunerate})}$ (same stem)
 - $\vec{w}^{(\text{iteration})}$, $\vec{w}^{(\text{consideration})}$... (same suffix $\approx$ same POS)

FASTTEXT (2) I

Assume, we want to represent the word *example*:

- Character n-grams ($n = 3$):

<ex, exa, xam, amp, mpl, ple, le>, <example>

- In practice, we don't set $n = a$ but rather $a \leq n \leq b$
- Character n-grams ($2 \leq n \leq 4$):

<e, ex, xa, am, mp, pl, le, e>,
<ex, exa, xam, amp, mpl, ple, le>,
<exa, exam, xamp, ampl, mple, ple>,
<example>

- Note, that the 4-gram *exam* is different from the word <exam>.

FASTTEXT (3) I

Representation of a known word:

Average of the word's embedding and char-n-gram embeddings

$$\vec{w}^{(i)} = \frac{1}{|\text{ngrams}(i)| + 1} \left[\vec{u}^{(i)} + \sum_{n \in \text{ngrams}(i)} \vec{u}^{(n)} \right]$$

Representation of an unknown word:

Average of char-n-gram embeddings

$$\vec{w}^{(i)} = \frac{1}{|\text{ngrams}(i)|} \sum_{n \in \text{ngrams}(i)} \vec{u}^{(n)}$$

FASTTEXT TRAINING I

- ngrams typically contains character 3- to 6-grams
- Replace $\vec{w}$ in Skipgram objective with its new definition
- During backpropagation, loss gradient vector $\frac{\partial L}{\partial \vec{w}^{(i)}}$ is distributed to word vector $\vec{u}^{(i)}$ and associated n-gram vectors $\vec{u}^{(n)}$

SUMMARY I

- Word2Vec as a bigram Language Model
- Negative Sampling
- Skipgram: Predict words in window given word in the middle
- CBOW: Predict word in the middle given words in window
- fastText: N-gram embeddings generalize to unseen words
- Any questions?

USING PRETRAINED EMBEDDINGS I

- Knowledge *transfer* from unlabelled corpus
- Design choice: Fine-tune embeddings on task or freeze them?
 - Pro: Can learn/strengthen features that are important for task
 - Contra: Training vocabulary is small subset of entire vocabulary → we might overfit and mess up topology w.r.t. unseen words

INITIALIZING NN WITH PRETRAINED EMBEDDINGS I

Model	MR	SST-1	SST-2	Subj	TREC	CR	MPQA
CNN-rand (randomly initialized)	76.1	45.0	82.7	89.6	91.2	79.8	83.4
CNN-static (pretrained+frozen)	81.0	45.5	86.8	93.0	92.8	84.7	89.6
CNN-non-static (pretrained+fine-tuned)	81.5	48.0	87.2	93.4	93.6	84.3	89.5
CNN-multichannel (combination)	81.1	47.4	88.1	93.2	92.2	85.0	89.4

► Kim (2014): Convolutional Neural Networks for Sentence Classification

RESOURCES I

- <https://fasttext.cc/docs/en/crawl-vectors.html>
 - Embeddings for 157 languages, trained on big web crawls, up to 2M words per language
- <https://nlp.stanford.edu/projects/glove/>
 - GloVe word vectors: Co-occurrence-count objective, not n-gram based

ANALOGY MINING (1) I

country-capital

$$\vec{w}(\text{Tokio}) - \vec{w}(\text{Japan}) + \vec{w}(\text{Poland}) \approx \vec{w}(\text{Warsaw})$$


opposite

$$\vec{w}(\text{unacceptable}) - \vec{w}(\text{acceptable}) + \vec{w}(\text{logical}) \approx \vec{w}(\text{illogical})$$

Nationality-adjective

$$\vec{w}(\text{Australian}) - \vec{w}(\text{Australia}) + \vec{w}(\text{Switzerland}) \approx \vec{w}(\text{Swiss})$$

ANALOGY MINING (2) I



ANALOGY MINING (3) I

$$\vec{w}^{(a)} - \vec{w}^{(b)} + \vec{w}^{(c)} = \vec{w}^{(?)}$$

$$\vec{w}^{(d)} = \operatorname{argmax}_{\vec{w}^{(d')} \in \mathbf{W}} \cos(\vec{w}^{(?)}, \vec{w}^{(d')})$$

Table 8: *Examples of the word pair relationships, using the best word vectors from Table 4 (Skip-gram model trained on 783M words with 300 dimensionality).*

Relationship	Example 1	Example 2	Example 3
France - Paris	Italy: Rome	Japan: Tokyo	Florida: Tallahassee
big - bigger	small: larger	cold: colder	quick: quicker
Miami - Florida	Baltimore: Maryland	Dallas: Texas	Kona: Hawaii
Einstein - scientist	Messi: midfielder	Mozart: violinist	Picasso: painter
Sarkozy - France	Berlusconi: Italy	Merkel: Germany	Koizumi: Japan
copper - Cu	zinc: Zn	gold: Au	uranium: plutonium
Berlusconi - Silvio	Sarkozy: Nicolas	Putin: Medvedev	Obama: Barack
Microsoft - Windows	Google: Android	IBM: Linux	Apple: iPhone
Microsoft - Ballmer	Google: Yahoo	IBM: McNealy	Apple: Jobs
Japan - sushi	Germany: bratwurst	France: tapas	USA: pizza

SUMMARY I

- Applications of Word Embeddings
 - Word vector initialization in neural networks for NLP tasks
 - E.g., sentiment classification of reviews, topical classification of news
 - Analogy mining
 - Information retrieval: semantic search, query expansion
 - Simple and fast aggregations of sentence representations
 - ...
- Any questions...?