


Deep Learning for NLP

Transformer

Long Sequences: Transformer-XL



Learning goals

- Understand the limitations for long sequences
- Understand the Segment Recurrence mechanism
- Understand relative positional encodings

LIMITATION OF THE TRANSFORMER I

Table 1: Maximum path lengths, per-layer complexity and minimum number of sequential operations for different layer types. n is the sequence length, d is the representation dimension, k is the kernel size of convolutions and r the size of the neighborhood in restricted self-attention.

Layer Type	Complexity per Layer	Sequential Operations	Maximum Path Length
Self-Attention	$O(n^2 \cdot d)$	$O(1)$	$O(1)$
Recurrent	$O(n \cdot d^2)$	$O(n)$	$O(n)$
Convolutional	$O(k \cdot n \cdot d^2)$	$O(1)$	$O(\log_k(n))$
Self-Attention (restricted)	$O(r \cdot n \cdot d)$	$O(1)$	$O(n/r)$

not cool

cool

Source: Vaswani et al. (2017)

Advantage:

- Every token can *directly* attend to each other token
- Cf. RNN: At worst n sequential operations (last to first token)

Severe Limitation:

- Every token attends to each other token (incl. itself)
→ We need to calculate n^2 attention weights
- Computational complexity of Transformer scales quadratically with the sequence length
→ Longer sequences are disproportionately expensive

TRANSFORMER-XL I

Key facts:

- Objective: Autoregressive Language Modeling task
- Transformer decoder model
- Addresses long sequences
- Assumption: *No* infinite memory & compute; limited resources
- (Possible) Solution Vanilla Transformer:
 - Split corpus into shorter segments
 - Limited contextual information
- Solution Transformer-XL:
 - Segment-level recurrence mechanism
 - Able to model longer-term dependencies

TRANSFORMER-XL I




Figure 1: Illustration of the vanilla model with a segment length 4.

Source: Dai et al. (2019)

- Contextual information limited to segments
- Does not respect semantic or syntactic boundaries

TRANSFORMER-XL I




Figure 2: Illustration of the Transformer-XL model with a segment length 4.

Source: Dai et al. (2019)

- Caches hidden states from the previous segment
- Contextual information flows across segments

SEGMENT RECURRENCE I

- Let $s_\tau = [x_{\tau,1}, \dots, x_{\tau,L}]$ and $s_{\tau+1} = [x_{\tau+1,1}, \dots, x_{\tau+1,L}]$ be two consecutive segments of length L .
- Let $h_\tau^n \in \mathbb{R}^{L \times d}$ denote the n -th layer hidden states for s_τ .
- Using segment recurrence, the n -th layer hidden states for the following segment $s_{\tau+1}$ are computed as follows:

$$\tilde{h}_{\tau+1}^{n-1} = \text{Concat}[SG(h_\tau^{n-1}), h_{\tau+1}^{n-1}]$$

$$q_{\tau+1}^n = h_{\tau+1}^{n-1} W_q^T; \quad k_{\tau+1}^n = \tilde{h}_{\tau+1}^{n-1} W_k^T; \quad v_{\tau+1}^n = \tilde{h}_{\tau+1}^{n-1} W_v^T$$

$$h_{\tau+1}^n = \text{Trafo}(q_{\tau+1}^n, k_{\tau+1}^n, v_{\tau+1}^n),$$

where $SG(\cdot)$ stands for "stop-gradient".

RELATIVE POSITIONAL ENCODINGS I

Problem:

- *Absolute* positional encodings (PE) would assign the same embedding to words in similar positions in both segments
- No positional difference between $x_{\tau,j}$ and $x_{\tau+1,j}$

Solution:

- Inject information about the relative distance between a query vector and the respective key vectors directly into the Attention mechanism
- *Comment:* Using relative PEs utterly necessary here, but also applicable independently of the segment recurrence

RELATIVE POSITIONAL ENCODINGS II

$$\begin{aligned}\mathbf{A}_{i,j}^{\text{abs}} &= \underbrace{\mathbf{E}_{x_i}^\top \mathbf{W}_q^\top \mathbf{W}_k \mathbf{E}_{x_j}}_{(a)} + \underbrace{\mathbf{E}_{x_i}^\top \mathbf{W}_q^\top \mathbf{W}_k \mathbf{U}_j}_{(b)} \\ &+ \underbrace{\mathbf{U}_i^\top \mathbf{W}_q^\top \mathbf{W}_k \mathbf{E}_{x_j}}_{(c)} + \underbrace{\mathbf{U}_i^\top \mathbf{W}_q^\top \mathbf{W}_k \mathbf{U}_j}_{(d)}.\end{aligned}$$

$$\begin{aligned}\mathbf{A}_{i,j}^{\text{rel}} &= \underbrace{\mathbf{E}_{x_i}^\top \mathbf{W}_q^\top \mathbf{W}_{k,E} \mathbf{E}_{x_j}}_{(a)} + \underbrace{\mathbf{E}_{x_i}^\top \mathbf{W}_q^\top \mathbf{W}_{k,R} \mathbf{R}_{i-j}}_{(b)} \\ &+ \underbrace{\mathbf{U}_i^\top \mathbf{W}_{k,E} \mathbf{E}_{x_j}}_{(c)} + \underbrace{\mathbf{V}_i^\top \mathbf{W}_{k,R} \mathbf{R}_{i-j}}_{(d)}.\end{aligned}$$

RELATIVE POSITIONAL ENCODINGS I

Solution:

- Replace all absolute PEs with relative ones (fixed + sinusoidal)
→ $\mathbf{R}_{i-j}$ instead of $\mathbf{U}_j$ in (b) and (d)
- Replace all positional query vectors with single trainable embeddings
→ u and v instead of $\mathbf{U}_i^\top \mathbf{W}_q^\top$ in (c) and (d)
- Use separate weight matrices for linearly projecting $\mathbf{E}$
→ $\mathbf{W}_{k,E}$ and $\mathbf{W}_{k,R}$ instead of $\mathbf{W}_k$