


Deep Learning for NLP

Post-BERT Era Model Distillation



Learning goals

- Soft vs. hard targets
- Understand how distillation works
- DistilBERT
- Other approaches towards compression

MODEL COMPRESSION

Motivation: ► Bucila et al., 2006

- Existence of computationally expensive, cumbersome ensemble models
- Accuracy/Performance not everything that should be taken into account (also interpretability and deployability)
- Time and space requirements also of importance
- *Model Compression* aims “to obtain fast, compact yet highly accurate models”
- The fast and compact model should approximate the function learned by the slower and larger model

MODEL COMPRESSION

Advantages:


- Compressing the knowlegde of the ensemble into a single model has the benefits of
 - easier deployment
 - better generalization
 - (potential) interpretability
- *Reasoning:*
 - Cumbersome model generalizes well, because it is the average of an ensemble
 - Small model trained to generalize in the same way typically better than small model trained "the normal way"

MODEL COMPRESSION

- *Challenge:*
 - Bucila et al. assumed only knowledge about the weights of the large (ensemble) model
 - No access to its training data
 - Their work is mainly on generating suitable pseudo data for model compression
- *NLP:*
 - Abundant amounts of data; no need for creating pseudo data
 - Self-supervised objectives can be used for model compression

MODEL DISTILLATION

Motivation:



► Source: Sanh et al., 2019

MODEL DISTILLATION

Types of targets

- “Hard” Targets

$[0, 0, 0, 1, 0]$

- Present when learning from labeled data
- Ordinary way of training models
- Knowledge transfer possible via “soft” targets from original model
- Possible “soft” targets
 - Logits + L2 loss ► Bucila et al., 2006
 - Softmax scores
 - Temperature T in the softmax ► Hinton et al., 2015

$$q_i = \frac{\exp(z_i/T)}{\sum_j \exp(z_j/T)}$$

MODEL DISTILLATION

Temperature:

$$q_i = \frac{\exp(z_i/T)}{\sum_j \exp(z_j/T)}$$

- *Why?* Produce a softer probability distribution over classes
- z_i/z_j are the raw logits
- Parameter T controls to smoothness of the distribution
 - Higher T : “softer” distribution
 - Lower T : more pronounced distribution
- *Distillation phase*: Apply same temperature to teacher and student
- *Inference*: Set $T = 1$ to recover standard softmax

MODEL DISTILLATION

Intuition

- Train the student to *mimick* the teacher's behaviour
- **Question:** Advantage of the soft targets:
 - No need to know the true labels of the training instances
 - More fine-grained information about the teacher's behaviour
- When true labels are known:
 - Weighted average of two different objective functions possible
 - On soft targets (*Be close to the teacher*) and
 - on hard targets (*be close to the "ground truth"*)

DISTILBERT

Characteristics of the student architecture:

- Half the number of layers compared to BERT¹ (6 instead of 12)
- Remove segment embeddings² (required for NSP objective)
- Half of the size of BERT, but retains up to 97% of the performance on GLUE
- Initialize from BERT (taking one out of two hidden layers)
- Same pre-training data as BERT (Wiki + BooksCorpus)

¹ Rationale for "only" reducing the number of layers:

Larger influence on the computation efficiency compared to e.g. hidden size dimension

² Sanh et al. call them "token-type embeddings" in their paper

DISTILBERT

Combination of three different loss functions:

- *Soft targets:*

Distillation loss $L_{ce} = \sum_i t_i \cdot \log(s_i)$

- *Hard targets:*

Masked language modeling loss L_{mlm} (cf. BERT)

- *Alignment:*

Cosine-Embedding-Loss $L_{cos} = 1 - \frac{\vec{w}^{(i)T} \vec{w}^{(j)}}{\|\vec{w}^{(i)}\|_2 \cdot \|\vec{w}^{(j)}\|_2}$

(*Rationale:* Keep DistilBERT's embeddings close to BERT's)

Adopt improvement from RoBERTa:

- Stop using the NSP loss (hence removed segment embeddings)
- Dynamic masking for MLM
- Train with large batches

DISTILBERT

Performance differences to BERT:

Table 1: **DistilBERT retains 97% of BERT performance.** Comparison on the dev sets of the GLUE benchmark. ELMo results as reported by the authors. BERT and DistilBERT results are the medians of 5 runs with different seeds.

Model	Score	CoLA	MNLI	MRPC	QNLI	QQP	RTE	SST-2	STS-B	WNLI
ELMo	68.7	44.1	68.6	76.6	71.1	86.2	53.4	91.5	70.4	56.3
BERT-base	79.5	56.3	86.7	88.6	91.8	89.6	69.3	92.7	89.0	53.5
DistilBERT	77.0	51.3	82.2	87.5	89.2	88.5	59.9	91.3	86.9	56.3

► Source: Sanh et al., 2019

Takeaways:

- Largest performance on textual entailment task (RTE)
- Even better reported performance on Winograd schemas (WNLI)
- *Overall:* Very consistent performance a little worse than BERT

DISTILBERT

Ablation study regarding the loss:

Table 4: **Ablation study.** Variations are relative to the model trained with triple loss and teacher weights initialization.

Ablation	Variation on GLUE macro-score
$\emptyset - L_{cos} - L_{mlm}$	-2.96
$L_{ce} - \emptyset - L_{mlm}$	-1.46
$L_{ce} - L_{cos} - \emptyset$	-0.31
Triple loss + random weights initialization	-3.69

► Source: Sanh et al., 2019

Takeaways:

- Start training from BERT's weights is crucial
- Removing MLM loss has little impact
 - Signal in the soft targets sufficiently strong
- Notable contribution of the cosine embedding loss

DISTILBERT

Size and speed:

Table 3: **DistilBERT is significantly smaller while being constantly faster.** Inference time of a full pass of GLUE task STS-B (sentiment analysis) on CPU with a batch size of 1.

Model	# param. (Millions)	Inf. time (seconds)
ELMo	180	895
BERT-base	110	668
DistilBERT	66	410

► Source: Sanh et al., 2019

RELATED APPROACHES (1)

Quantization:

- *In General*: Mathematical procedure of converting a number from a continuous scale to a discrete scale
- *In Machine Learning (usually)*: Convert a number from high precision (float16/float32) to a lower precision (float8/int8)
- For models with billions of parameters, this can save enormous amount of memory
- Given that model sizes (at the moment) grow faster than GPU memory, this can be crucial for even making *inference* of large models feasible
- *Trade-Off*: Memory savings vs. model quality

RELATED APPROACHES (2)

Pruning:

- Remove (“*prune*”) certain parts of the model according to some score measuring the importance of that part for performance
- Common importance score (in NLP):
Sensitivity of the loss wrt the values of the neurons

► Yang et al., 2022

$$\text{IS}(\Theta) = \mathbb{E}_{x \sim \mathcal{X}} \left| \frac{\partial \mathcal{L}(x)}{\partial \Theta} \Theta \right|$$

- Common choices for Θ (in Transformers):
 - Output of an attention head
 - An intermediate neuron in the FFN layer

RELATED APPROACHES (2)

Self-supervised Pruning ► Yang et al., 2022

- Substitute $\mathcal{L}$ by

$$\mathcal{L}_{\text{KL}}(x) = \text{KL}(\text{stopgrad}(q(x)) || p(x)),$$

- where
 - $q(x)$ is the original model prediction distribution,
 - $p(x)$ is the to-be-pruned model prediction distribution,
 - and `stopgrad` is used to stop back-propagating gradients

SUMMARY

- Model compression/distillation against growing model sizes
- DistilBERT as a powerful, small alternative to BERT
- Other approaches (pruning) that allow for a more targeted “deactivation” of individual heads (and might thus also improve interpretability)
- Quantization as an approach to *significantly* reduce the memory requirements of large models