


Deep Learning for NLP

Training LLMs

How to reduce compute and memory




Learning goals

- Learn about different techniques to reduce compute and memory
- Learn about distributed training with data/tensor parallelism
- Learn about FlashAttention

DISTRIBUTED TRAINING

- Training LLMs on several GPUs is faster than on one, but we need to figure out how to distribute.
- Avoiding OOM (out of memory) issues
- **Data parallelism:** split the data on different model replicas
- **Tensor parallelism:** split model parameters accross GPUs

DATA PARALLELISM (1) (ANIMATED GIF!)



Source: Nvidia

DATA PARALLELISM (2)

- **Data Splitting**

- Dataset divided into smaller chunks. Each chunk assigned to a different GPU.
- Each node processes a different subset of the data in parallel.

- **Model Replication**

- Each GPU has a replica of the neural network model
- These replicas are trained independently on their respective data subsets

- **Gradient Aggregation (All Reduce)**

- Gradients are computed locally on each node and then averaged across nodes.

- **Parameter Synchronization**

- Model parameters are updated synchronously across nodes.
- This ensures that all model replicas remain consistent with each other after each update step.

TENSOR PARALLELISM (2)

- **Model Partitioning:**

- The model's layers or tensors are split across multiple devices
- Different parts of the model are assigned to different devices, enabling them to work on separate portions of the computations simultaneously

- **Forward and Backward Passes:**

- During the forward pass, each device processes its portion of the tensors with intermediate results passed between devices
- In the backward pass, gradients are computed in the reverse order, again with necessary data transfers between devices

- **Parameter Updates:**



- Parameter updates can be performed independently on each device for the parameters they own
- After each update step, the devices synchronize to ensure consistency across the distributed model

FlashAttention

Fast and Memory-Efficient Exact Attention with IO-Awareness

- Memory-efficient
 - Reducing memory requirements from $O(n^2)$ to $O(n)$
Question: Why quadratic?
- IO aware
 - Reducing memory load/store operations
- Therefore: much faster than the regular implementation
 - The bigger the attention matrices, the more you gain (until $O(n)$ no longer fits).
- Exact
 - Same as “vanilla attention”, not an approximation

GPU MEMORY HIERARCHY



► Source: Dao, 2022



COMPUTING CONSIDERATIONS

- GPU computations are extremely fast, but data transfer is slow.
- Typical GPU computations are elementwise operations with a huge number of elements, resulting in high memory access
- → GPU has to wait for data.
- → Transformer operations are memory-bound.
- IO aware means reducing memory load/store operations

FLASH ATTENTION


- Tiling or chunking of matrices into blocks
- Operation fusion
 - For example: first matrix multiplication (query-key), softmax, second matrix multiplication (weight-value)
- This minimizes redundant swapping in and out of data
- General idea: you avoid that the large intermediate tensors (e.g., pre-softmax) are written to HBM and then read back into SRAM memory

SRAM VS HBM / FUSED KERNEL




► Source: Dao, 2022

TILING



OPERATION FUSION



Source: https://horace.io/brrr_intro.html

LIMITATIONS AND PROSPECTS

- FlashAttention requires implementing “custom” attention in CUDA
 - A new CUDA kernel for each new attention implementation
 - CUDA is lower-level than PyTorch
 - Implementation not transferable accross GPUs
- Towards IO-Aware Deep Learning
 - Extending beyond attention
 - DeepSeek

TEACHING EVALUATION

► teaching evaluation: QAEQ7

