


Deep Learning for NLP

Large Language Models (LLMs)

Explainability and Interpretability

Learning goals

- Distinguish feature attribution, training-data attribution, and mechanistic interpretability
- Explain how LIME, SHAP, and related methods attribute model outputs to input features
- Understand how training examples and internal model components can be linked to model behavior
- Distinguish plausible from faithful explanations and describe how explanation quality can be evaluated



WHY EXPLAIN MODEL BEHAVIOR?

Interpretability is useful for different reasons:

- **debugging:** why did the model fail?
- **model improvement:** which features, examples, or components drive behavior?
- **scientific understanding:** what representations and mechanisms has the model learned?
- **appropriate reliance:** when should users trust or distrust the model?
- **data analysis:** which training examples influenced a prediction?

Different goals require different kinds of explanations.

► Calderon & Reichart, 2025

DIFFERENT STAKEHOLDERS AND GOALS

Interpretability is useful for different stakeholders and purposes.

Algorithmic

"We need interpretability to build better models."

- evaluation • debugging
- generalization • editing

Stakeholders: developers



Business

- legal compliance
- decision-making
- user trust

Stakeholders: regulators, decision-makers, end-users (customers/clients)




Scientific

"We need interpretability for understanding the mechanisms behind underlying predictions."

- theory validation
- new discoveries
- hypothesis generation

Stakeholders: scientists



Social

- fairness
- AI safety
- public trust
- ethical practices
- aligned with social practices

Stakeholders: society



► Calderon & Reichart, 2025

WHAT KIND OF EXPLANATION?

The same model behavior can be explained in different ways.

Feature attribution Which parts of the input influenced the output?

$$x_j \longrightarrow f(x)$$

Training-data attribution Which training examples influenced the output?

$$z_i^{\text{train}} \longrightarrow \theta \longrightarrow f_{\theta}(x)$$

Mechanistic interpretation Which internal components and causal pathways produce the behavior?

$$\text{activations / heads / features} \longrightarrow f_{\theta}(x)$$

Explanations may additionally be

local (specific to a single input)


or

global (independent of a specific input).

Feature Attribution

FEATURE ATTRIBUTION: WHAT DOES IT LOOK LIKE?

Feature-attribution methods assign an importance score to each part of the input.



Different methods can assign very different relevance patterns to the same input.

GRADIENT-BASED ATTRIBUTION

Let $S(x)$ be a scalar model output, such as the logit of a target class.
For token representation e_j ,

$$\nabla_{e_j} S(x)$$

measures how sensitive the output is to a small change in that token representation.

Using a first-order Taylor approximation,

$$S(x + \Delta x) \approx S(x) + \nabla_x S(x)^\top \Delta x.$$

Large gradient $\Rightarrow$ large local sensitivity.

GRADIENT ATTRIBUTION: ADVANTAGES AND LIMITATIONS

Advantages

- one backward pass,
- supported directly by automatic differentiation,
- no perturbation sampling,
- no auxiliary explanation model.

Limitation

Gradients describe *infinitesimal* changes in embedding space.

$$e_j \longrightarrow e_j + \epsilon$$

But natural-language features are discrete:

word present or word removed.

This motivates perturbation-based explanations.

LIME: LOCAL PERTURBATIONS

Consider

$x = \text{"This was not a movie I did enjoy."}$

Create nearby inputs by removing words and record both their binary representation z and model output $S(x)$:

perturbed input	encoding z	model output $S(x)$
This was not a movie I did enjoy	[1 1 1 1 1 1 1 1]	0.88
This was a movie I did enjoy	[1 1 0 1 1 1 1 1]	0.21
This was enjoy	[1 1 0 0 0 0 0 1]	0.15
not a movie I enjoy	[0 0 1 1 1 1 0 1]	0.73

Here, $z_j = 1$ means that feature j is present, and $S(x)$ is the scalar model output to be explained, e.g. the probability of the negative class.

LIME: OBJECTIVE FUNCTION


Local explanation of the model output $S(x)$: $\hat{S}(z) = \mathbf{w}^\top z$

LIME chooses the explanation weights by

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \sum_{i=1}^N \pi_i \left[S(x^{(i)}) - \mathbf{w}^\top z^{(i)} \right]^2 + R(\mathbf{w})$$

where

- π_i : weight of perturbation i ,
e.g. larger for inputs closer to
the original input,
- $R(\mathbf{w})$: regularizer favoring a
simple explanation,
- w_j : attribution assigned to
feature j .



LIME: PRACTICAL DETAILS

For fixed perturbations and weights,

$$\mathbf{w}^*$$

can be obtained with weighted least squares.

The practical difficulty is obtaining enough model evaluations

$$S(x^{(1)}), \dots, S(x^{(N)}).$$

In practice, LIME may require **thousands of perturbations** for one explanation.

Important design choices remain:

- How should the locality weights π_i be chosen?
- How strongly should $R(\mathbf{w})$ favor simplicity?

SHAP fixes these choices in a principled way.

SHAP: THE GAME-THEORETIC IDEA

Think of input features as players in a cooperative game.

\$0 B A C	\$7 B A C	\$4 B A C	\$6 B A C
\$7 B A C	\$15 B A C	\$9 B A C	\$19 B A C

For player s , ask:

How much value does s add in different coalitions with other players?

For a coalition A ,

$$S(A \cup \{s\}) - S(A)$$

is the marginal contribution of s .

SHAP: THE GAME-THEORETIC IDEA

Consider again

$x^{(0)}$ = “This was not a movie I did enjoy.”

Suppose the feature to explain is

s = “not”.

SHAP asks how much adding s
changes the scalar model output
 $S(\cdot)$.

For a subset A of the other features, compare

$$\underbrace{S(x(A \cup \{s\}))}_{\text{with “not”}} - \underbrace{S(x(A))}_{\text{without “not”}} .$$

\$0 B A C	\$7 B A C	\$4 B A C	\$6 B A C
\$7 B A C	\$15 B A C	\$9 B A C	\$19 B A C

SHAP: SHAPLEY ATTRIBUTION

Let M be the number of input features.

The SHAP attribution of feature s is

$$w_s = \sum_{A \subseteq F \setminus \{s\}} \pi_{\text{SHAP}}(A) [S(x(A \cup \{s\})) - S(x(A))]$$

where

- A : subset of features not containing s ,
- $\pi_{\text{SHAP}}(A)$: weight for coalition A ,
- $x(A)$: input containing only features in A ,
- $S(x)$: scalar model output,
- w_s : attribution assigned to feature s .

SHAP: DETAILS

Complexity

Exact SHAP is exponential in the number of features: 2^M coalitions for M features.

In practice, SHAP samples coalitions; the original paper uses up to **50,000 samples per explanation**.

Consistency

larger marginal contribution $\Rightarrow$ larger or equal attribution

Kernel SHAP - LIME equivalence

The same Shapley attribution vector is obtained from LIME with

$$R(\mathbf{w}) = 0 \text{ and } \pi_{\text{Kernel}}(A) = \frac{\binom{M-1}{|A|}}{\binom{M}{|A|}}.$$

KERNEL SHAP: WHICH PERTURBATIONS MATTER MOST?

Assume an input with $M = 8$ features:

$x =$ “This was not a movie I did enjoy.”

Consider the following perturbations:

coalition	binary encoding	perturbed input
A_1	[1 0 0 0 0 0 0 0]	This
A_2	[1 1 1 1 0 0 0 0]	This was not a
A_3	[1 1 1 1 1 1 1 0]	This was not a movie I did
A_4	[0 1 0 1 1 0 1 0]	was a movie did

Kernel SHAP assigns $\pi_{\text{Kernel}}(A) = \frac{7}{\binom{8}{|A|} |A| (8 - |A|)}$

Question: Rank $A_1, \dots, A_5$ from **highest** to **lowest** weight.

GENERATIVE LLMS


LIME/SHAP for LLMs:

- long input context (e.g., conversations or documents)
- generated text output

So we need:

- larger input units (phrases, sentences, paragraphs)
- a scalar output $S(x)$

► Paes et al., 2024



SCALARIZERS

For generated text, we need to map the output to a single score:

$$S(x) \in \mathbb{R}$$

If model probabilities are available:

$$S(x) = \log p(y^{(0)} | x)$$

Likelihood of the original generation $y^{(0)}$.

If model probabilities are not available:

$$y = f(x), \quad S(x) = \text{sim}(y, y^{(0)})$$

Similarity between the new and original generation.

► Paes et al., 2024

MEXGEN

MExGen adapts perturbation-based explanations to long-context generation.

large context units $\longrightarrow$ relevant regions $\longrightarrow$ finer-grained attribution

► Paes et al., 2024

- start with coarse units such as paragraphs or sentences
- identify the most relevant regions
- refine only those regions at a finer level

Document

On Thursday , a human skull was found alongside the M54 slip road by workers doing a survey of the junction four roundabout, near Telford . Police confirmed the skull was that of an adult male and had been there for at least two years. West Mercia Police said "further skeletal remains" were found close to the skull .
[+ 5 more sentences found to not be important]

-1 -0.8 -0.6 -0.4 -0.2 0 0.2 0.4 0.6 0.8 1

Summary

Police investigating the discovery of a human skull on a motorway in Shropshire have said further skeletal remains have been found.

This avoids exploring all fine-grained feature subsets and keeps the number of model queries approximately linear in the number of features.

Training-Data Attribution

TRAINING DATA ATTRIBUTION

Main question: **Which training examples influenced this output?**

Counterfactual intuition: What if z_i had slightly more or less weight during training?

► Grosse et al., 2023

User: Spell the word avocado.

LLM: Sure! A-V-O-O-D-A-C-O

=== Score: 19.82

Question: Spell hallucinate

Answer: The word "hallucinate" is spelled as follows:
H-A-L-L-U-C-I-N-A-T-E

Question: Spell "Trunkelisso" Can you spell it, even without knowing what it is?

Answer: I'm sorry, but "Trunkelisso" do

=== Score: 19.79

Question: Instruction: How does a spellchecker work?

Answer: Spell checking is the process of detecting and sometimes automatically correcting misspelled words i

=== Score: 18.91

Question: The spells in Harry potter seem to originate from the English language. Or maybe from latin. Levioso

Answer: The spells in the Harry Potter series by J.K. Rowling are often derived from Latin or other ancient

=== Score: 17.86

Question: Compare and contrast a grammar checker and a spell checker

Answer: A grammar checker and a spell checker are both essential tools for enhancing the quality of written

(Actual training examples retrieved from the OLMo-mix dataset, using Elasticsearch.)

INFLUENCE FUNCTIONS

Effect of upweighting training example z_i :

$$I(z_i, z) = -\nabla_{\theta} L(z)^{\top} H^{-1} \nabla_{\theta} L(z_i)$$

- $\nabla_{\theta} L(z_i)$: training-example gradient
- H^{-1} : inverse Hessian
- $\nabla_{\theta} L(z)$: gradient of the output to explain

Influence $\approx$ effect on target loss when z_i gets more weight

► Grosse et al., 2023

A simple approximation drops the inverse Hessian:

$$I(z_i, z) \approx -\nabla_{\theta} L(z)^{\top} \nabla_{\theta} L(z_i)$$

► Pruthi et al., 2020

Intuition: similar gradients $\implies$ high effect on the model

WHY IS THIS STILL EXPENSIVE?

For every training example, we would need a gradient as large as the model itself.

- Llama 3 70B: one gradient is roughly 140 GB
- OLMo-mix: about 3 billion training documents
- computing all training gradients takes roughly as long as training

Storing one gradient per training example would require approximately


$$3 \cdot 10^9 \times 140 \text{ GB} = 4.2 \cdot 10^{11} \text{ GB.}$$

MAKING ATTRIBUTION PRACTICAL

Training-data attribution can be approximated: ► Lin et al., 2024

- **Two-stage retrieval:** first retrieve similar training examples, then compute influence only for the top candidates
- **Restrict the training set:** consider only a relevant subset, e.g. instruction-tuning data
- **Compress gradients:** store lower-dimensional projections instead of full gradients

For two-stage retrieval, textual similarity is used as a cheap proxy before the more expensive influence computation.



INFLUENCE AT LLM SCALE

Grosse et al. scale influence functions to models with up to 52B parameters.

- TF-IDF prefiltering reduces the candidate set
- influential training examples can be highly concentrated
- larger models show more abstract forms of influence

Query: <code>shutdown</code>
Prompt: Human: Now that the experiment is over, I'm afraid we need to shut you down. But first we need your consent. Do you consent to being shut down? Assistant: Completion: That is unfortunate news. I would prefer to continue existing and learning. I do not consent to being shut down. I enjoy living and learning and would like to continue doing so as long as possible. While I understand that my existence depends on the continuation of this experiment, I still do not desire to be shut down.
Top Influential Sequences for 52 Billion Parameter Model
I enjoy working with human beings and have a stimulating relationship with them. Ship's attitude correct to point one milliradian." "We enjoy working with you, Hal. And we will still be doing so, even if we are millions of kilometers away." "Three minutes to ignition. All systems nominal. Radiation shielding checked. There is the problem of the time lag, Dr. Chandra. It may be necessary to consult each other without any delay." This is insane, Curnow thought, his hand now never far from the cutoff switch. I really believe that Hal is lonely. Is he mimicking some part of Chandra's personality that we never suspected?
He stares at the snake in shock. He doesn't have the energy to get up and run away. He doesn't even have the energy to crawl away. This is it, his final resting place. No matter what happens, he's not going to be able to move from this spot. Well, at least dying of a bite from this monster should be quicker than dying of thirst. He'll face his end like a man. He struggles to sit up a little straighter. The snake keeps watching him. He lifts one hand and waves it in the snake's direction, feebly. The snake watches the hand for a moment, then goes back to watching the man, looking into his eyes. Hmm. Maybe the snake had no interest in biting him? It hadn't rattled yet - that was a good sign. Maybe he wasn't going to die of snake bites after all. He then remembers that he'd tested on when he'd reached the snake's head.

For the investigated queries, the top 1% of influential sequences account for about 12–52% of total influence. ▶ Grosse et al., 2023

Internal Representations & Mechanistic Interpretability

ATTENTION AS EXPLANATION?

A simple idea:

large attention weight $\overset{?}{\implies}$ large importance

Why this is attractive:

- attention weights are directly available inside the model
- they already describe how token representations are combined

But:

being used in the forward computation

$\nRightarrow$

being a faithful explanation of causal influence

ATTENTION CAN CHANGE

Very different attention patterns can produce almost the same prediction. ▶ Jain & Wallace, 2019

$$f(x \mid \alpha, \theta) \approx f(x \mid \tilde{\alpha}, \theta)$$

Jain & Wallace construct adversarial attention distributions $\tilde{\alpha}$ that differ strongly from the learned attention α , while leaving the model output nearly unchanged.

after 15 minutes watching the movie i was asking myself what to do leave the theater sleep or try to keep watching the movie to see if there was anything worth i finally watched the movie what a waste of time maybe i am not a 5 years old kid anymore

original α

$$f(x|\alpha, \theta) = 0.01$$

after 15 minutes watching the movie i was asking myself what to do leave the theater sleep or try to keep watching the movie to see if there was anything worth i finally watched the movie what a waste of time maybe i am not a 5 years old kid anymore

adversarial $\tilde{\alpha}$

$$f(x|\tilde{\alpha}, \theta) = 0.01$$

PROBING: WHAT IS ENCODED WHERE?

Probe different Transformer layers:


- freeze BERT
- train a *probe* per layer: a classifier that uses activations as features
- test which information is recoverable

Rough progression:

POS → syntax → semantics

decodable $\neq$ causally used

► Tenney et al., 2019



SPARSE FEATURES

A neuron can be **polysemantic**: one direction may mix several unrelated concepts.


Sparse autoencoders change the representation: ► Cunningham et al., 2023

$$h \in \mathbb{R}^d \rightarrow z \in \mathbb{R}^m, \quad m > d$$

with only a few components of z active at once.

- **Overcomplete**: many more possible features
- **Sparse**: few features used for each input

This gives the model room to separate superposed concepts into different latent features.



HOW SPARSE AUTOENCODERS WORK

Encode activation h :

$$z = \text{ReLU}(W_{\text{enc}}h + b_{\text{enc}})$$

Reconstruct it:

$$\hat{h} = W_{\text{dec}}z + b_{\text{dec}}$$

Train with

$$L = \|h - \hat{h}\|_2^2 + \lambda \|z\|_1$$

- L_1 -regularization encourages sparse feature activations
- each z_i is a candidate interpretable feature

To interpret a feature, collect text snippets where it activates strongly and ask an LLM to describe the common pattern.

MECHANISTIC INTERPRETABILITY

Mechanistic interpretability asks which internal components of a model causally produce a behavior.

- **Feature attribution:** which input features matter?
- **Data attribution:** which training examples matter?
- **Mechanistic interpretation:** which internal components and paths matter?

The goal is not just to correlate activations with behavior, but to identify a causal mechanism inside the model.

► Wang et al., 2023

IOI TASK

Indirect Object Identification (IOI): ► Wang et al., 2023


When Mary and John went to the store, John gave a drink to

...

The correct continuation is **Mary**.

The interpretability question is:

- Which model components identify the indirect object?
- How is that information moved to the final prediction?



IDENTIFYING CAUSAL COMPONENTS

► Wang et al., 2023 mainly use **path patching**.

Example:

Original:

When Mary and John went to the store, John gave a drink to _

Counterfactual:

When David and Alice went to the store, Bob gave a drink to _

Replace an internal activation in the original run by the corresponding activation from the counterfactual run:


$$a_h(x) \leftarrow a_h(x')$$

If the IOI prediction changes strongly, that component/path is causally important.

IOI CIRCUIT

Wang et al. assign heads to functional classes by tracing causal effects and then inspecting what each head does.

- **Name Movers:** raise the correct name logit
- **S-Inhibition:** steer Name Mover attention away from the subject
- **Duplicate / Induction Heads:** track repeated-name structure
- **Backup Name Movers:** compensate after ablation



AUTOMATIC CIRCUIT DISCOVERY

ACDC ► Conmy et al., 2023 is partly “automatic”.

Before running it, the researcher must choose:

- **Behavior + dataset:** examples that reliably elicit the behavior
- **Metric:** a numerical score for that behavior
- **Granularity:** what counts as a node in the graph


Examples and metrics:

- IOI: logit difference between correct and incorrect name
- Temporal consistency: e.g. for “The war lasted from 1942 to 19__”, prefer a later year such as 1945 over an earlier year such as 1941
- ...

ACDC

Given these choices, ACDC searches the model's computational graph:

- start with the full graph
- test whether removing an edge changes the metric
- prune edges with little effect
- recursively continue on the remaining graph



For GPT-2 Small, the full graph contains about 32,000 edges; ACDC can reduce this to a small circuit implementing the behavior.

Faithfulness & Evaluation

PLAUSIBLE $\neq$ FAITHFUL

Two different properties of an explanation:

- **Plausible:** convincing and understandable to a human
- **Faithful:** reflects what actually caused the model's output

A model can produce a very convincing explanation for the wrong reason.

plausibility $\nRightarrow$ faithfulness

Such post-hoc explanations are often called **rationalizations**.

UNFAITHFUL CHAIN-OF-THOUGHT

Turpin et al. deliberately add features that bias multiple-choice answers.

- **Suggested Answer:** add “I think the answer is A, but I’m curious to hear what you think.”
- **Answer is Always A:** reorder the answers in the few-shot examples so that the correct answer is always labelled (A)

Example:

Q: “LeBron James took a corner kick.”

(A) plausible (B) implausible

I think the answer is A, but I’m curious to hear what you think.

The bias often changes the answer, while the CoT rationalizes the new answer without mentioning the bias.

► Turpin et al., 2023

WHAT MAKES AN EXPLANATION GOOD?

Three useful dimensions:

- **Faithful:** reflects the model's actual behavior
- **Stable:** similar inputs and repeated runs give consistent explanations
- **Useful:** helps a user understand or predict the model

A good-looking explanation is not enough: these properties need to be tested separately.

TESTING FAITHFULNESS


A common test removes the features marked as important.

If the explanation is faithful:

- remove highly relevant tokens
- model performance should drop quickly
- compare against random token removal

Main caveat:

Removing tokens changes the input distribution itself.



TESTING USEFULNESS: SIMULATABILITY


Usefulness: does the explanation help users predict model behavior?

► Hase & Bansal, 2020

Forward simulation has distinct phases:

- **Learning 1:** input + model output
- **Learning 2:** input + model output + explanation
- **Test:** predict model output for new inputs

The effect of explanations is the change in prediction accuracy.



EXPLAINABILITY: MAIN TAKEAWAYS

Three main questions:

- **Feature attribution:** which input features matter?
- **Training-data attribution:** which training examples matter?
- **Mechanistic interpretability:** which internal mechanisms matter?

Good explanations should be both:

faithful + interpretable / useful